



Università degli studi di Cagliari

Facoltà di Ingegneria e Architettura

Computer Engineering, Cybersecurity and Artificial Intelligence

Secure messaging application

Network Security

Giancarlo Pederiva, Nicola Scema

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Scenario</b>	<b>4</b>
2.1	Example User Flow . . . . .	4
<b>3</b>	<b>System Architecture</b>	<b>5</b>
3.1	Overview . . . . .	5
3.2	Components . . . . .	5
3.3	Diagram . . . . .	7
3.4	Communication Flow . . . . .	8
<b>4</b>	<b>Design and Implementation</b>	<b>10</b>
4.1	Key Management . . . . .	10
4.2	Session Key Derivation . . . . .	10
4.3	Message Encryption . . . . .	11
4.3.1	AES-GCM . . . . .	11
4.4	Broadcast Chat Functionality . . . . .	12
4.4.1	Joining a Broadcast . . . . .	12
4.4.2	Sending Broadcast Messages . . . . .	12
4.4.3	Receiving Broadcast Messages . . . . .	13
4.4.4	Leaving a Broadcast . . . . .	13
4.4.5	Security Analysis . . . . .	13
4.5	Libraries Used . . . . .	15
4.6	Code Structure . . . . .	15
4.7	Interface . . . . .	15
<b>5</b>	<b>Evaluation and Conclusions</b>	<b>17</b>
5.1	Initial Handshakes . . . . .	17
5.2	Private (client-to-client) messages . . . . .	18
5.2.1	Broadcast (group) messages and group key updates . . . . .	19
5.3	What Wireshark reveals and what it does not . . . . .	19
5.4	Limitations . . . . .	20
5.5	Possible Improvements . . . . .	20
5.6	Conclusion . . . . .	21

# 1 Introduction

In the context of digital communications, security has become a fundamental priority, especially in messaging applications where the transmission of confidential and sensitive data is becoming more widespread. The objective of this project is to develop a chat system that provides secure end-to-end communication, allowing two or more users to communicate with each other while ensuring the confidentiality, authenticity, and integrity of exchanged messages.

The idea behind our project consists of an instant messaging client that uses a central server to establish communications between the various users connected to the server. To ensure the security of communications, each client is equipped with a pair of persistent cryptographic keys and an ephemeral key, per session, used to establish an encrypted channel created with a customised handshake protocol.

The main objectives of the project are:

- Design and implement a chat client with an easy and intuitive, and accessible graphical interface capable of managing multiple users simultaneously.
- Securely manage user identities by generating and storing public/private key pairs.
- Realize a secure handshake protocol to negotiate temporary session keys and establish crypted end-to-end communications between users.
- Encrypt and decrypt messages exchanged using modern cryptographic algorithms (X25519 for key agreement, AES-GCM for message encryption).
- Avoid man-in-the-middle (MITM) attacks through the use of digital signatures and public key fingerprints.

## 2 Scenario

The system is designed to simulate secure instant messaging between two or more users connected over potentially insecure networks such as public Wi-Fi or untrusted LANs. Communication takes place through a central server that acts purely as a message relay and is considered untrusted, thus it cannot decrypt, store, or alter the content of exchanged messages.

This setup reflects realistic conditions in which users might be geographically distributed, connecting from different devices and locations, while still requiring strong guarantees of confidentiality, integrity, and authentication.

### 2.1 Example User Flow

1. User A opens the chat client, connects to the server via a TLS-secured WebSocket channel, and announces their presence.
2. User B connects to the same server and initiates a session with User A.
3. Both clients exchange their public identity keys and ephemeral session keys using a secure handshake protocol.
4. Once the session is established, messages are encrypted end-to-end before being sent to the server for relay.
5. Even if intercepted, the messages cannot be decrypted by the server or any third party.

This scenario models a secure peer-to-peer conversation mediated by an untrusted relay, a pattern often used in modern secure messaging applications.

In addition to one-to-one sessions, the system also supports a **broadcast chat mode**, where multiple clients share a common group state. A fresh epoch secret is distributed among members whenever someone joins or leaves, ensuring that only current participants can access group messages. Messages are encrypted end-to-end and relayed to all members by the server, which remains oblivious to their content. This enables secure multi-user conversations with forward secrecy and minimal trust in the server.

## 3 System Architecture

### 3.1 Overview

The proposed system is built upon a **client-server** architecture in which the server acts as a message relay rather than a trusted intermediary. This choice comes from the security goal of end-to-end encryption (E2EE): all encryption and decryption occur exclusively on the client side, ensuring that the server is unable to read, modify, or forge messages.

The architecture supports multiple clients connected simultaneously over **secure WebSocket connections** (`wss://`), which are themselves protected by TLS. TLS provides confidentiality and integrity at the transport layer, preventing network-level attackers from interfering with the WebSocket channel. However, TLS alone is not sufficient for our threat model, because it protects only the channel between a client and the server, not between two clients. For this reason, the application implements an application-layer encryption protocol on top of TLS.

Our threat model assumes that:

- The network may be insecure (e.g., public Wi-Fi with potential man-in-the-middle attackers).
- The server may be honest-but-curious or even compromised, and thus cannot be trusted with message contents.
- Attackers may attempt to impersonate other users, replay messages, or alter them in transit.

The design therefore, ensures that even if the server or the network is compromised, no plaintext communication or private keys are exposed.

### 3.2 Components

The system is composed of three main components, each with its own security responsibilities:

**Client:** The client represents the end user inside the system. While the user simply types and reads messages, the client transparently manages key generation, session establishment, and message encryption to ensure communication remains private and authenticated, hiding the complexity of cryptographic operations from the user. Client's main responsibilities include:

- A graphical user interface (GUI) for user interaction.
- Generation and secure storage of a persistent **Ed25519** identity key pair.
- On-demand generation of ephemeral **X25519** keys for each new session.
- Execution of the handshake protocol to authenticate peers and derive shared session keys.
- Encryption and decryption of all outgoing and incoming messages using **AES-256-GCM**.

**Server:** The server is designed to be minimal and unintrusive. Instead of interpreting or modifying the data it handles, its main purpose is to coordinate connections and deliver ciphertext messages between participants. This architecture keeps trust requirements low: the system remains secure even if the server is compromised. Its tasks are limited to:

- Maintaining a list of connected users and their advertised public keys.
- Acting solely as a relay for encrypted messages.
- Never storing or processing plaintext messages.
- Operating over TLS-secured WebSockets to protect against network-level interference.

In order to support TLS, the server requires a certificate-key-key pair. Rather than relying on external tools such as OpenSSL to generate and provision these files, our implementation creates them automatically at startup if none are present in the `certs/` directory. The procedure generates a 2048-bit RSA private key and a self-signed X.509 certificate, valid for ten years, with both `localhost` and `127.0.0.1` specified as subject alternative names (SANs). The resulting files `server.key` and `server.crt` are stored for reuse across sessions.

This design choice simplifies deployment in a development or testing environment: TLS is always available without manual configuration steps. Compared to manually invoking OpenSSL, our approach eliminates external dependencies, ensures consistency across different machines, and integrates directly into the server’s startup routine. While self-signed certificates naturally trigger verification warnings in clients, they are sufficient for localhost experimentation and classroom use. In a production setting, however, these certificates should be replaced with ones issued by a trusted Certificate Authority (CA).

**Cryptographic Primitives:** The cryptographic layer is the foundation of the system’s security. Instead of relying on a single algorithm, it combines multiple well-established primitives, each addressing a specific aspect of confidentiality, integrity, and authenticity. Together, they form a robust protocol that balances efficiency with modern security guarantees. The primitives used are:

- Ed25519 for identity authentication and public key verification.
- X25519 for elliptic-curve Diffie–Hellman (ECDH) key exchange.
- HKDF with SHA-256 for deriving independent symmetric keys from the shared secret.
- AES-256-GCM for authenticated encryption of messages.

### 3.3 Diagram

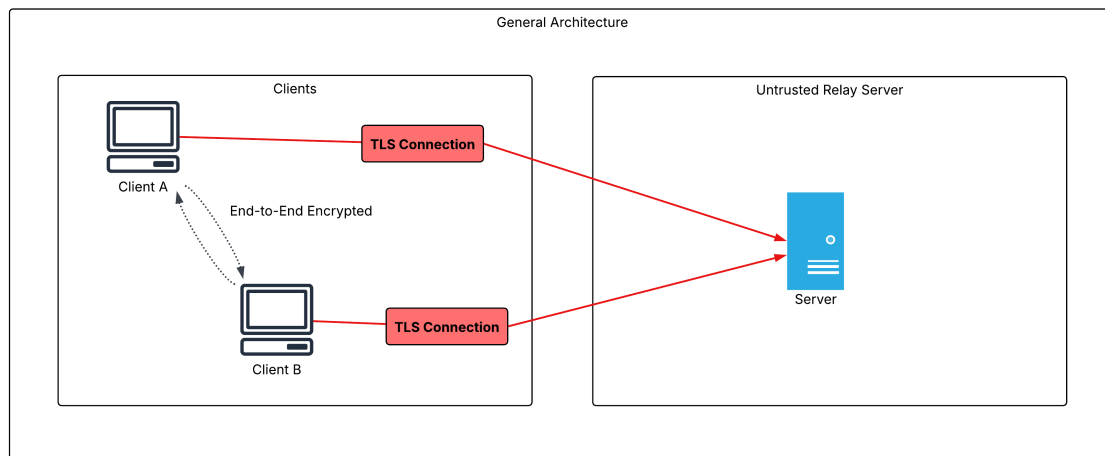


Figure 1: Logical system architecture showing secure end-to-end communication over an untrusted relay server.

In the diagram:

- Each client establishes a **TLS-protected WebSocket** connection to the relay server. This prevents passive eavesdropping and message tampering between the client and the server.
- During the handshake, clients exchange:
  1. Persistent identity keys (Ed25519) to authenticate themselves.
  2. Ephemeral Diffie–Hellman keys (X25519) to establish session-specific shared secrets.

These exchanges are signed to prevent man-in-the-middle attacks and verified against known key fingerprints.

- Once the handshake is complete, all messages are encrypted end-to-end using **AES-256-GCM** with keys derived from the ECDH output via HKDF.
- The server merely forwards the ciphertext without the ability to decrypt it.

This layered approach (TLS at the transport layer, E2EE at the application layer) ensures that:

1. Network attackers cannot intercept or modify WebSocket traffic.
2. The server cannot read or alter messages, even if compromised.
3. Compromise of one session’s keys does not endanger past or future sessions (**forward secrecy**).

### 3.4 Communication Flow

The high-level communication process works as follows:

1. Each client connects to the server over a TLS-secured WebSocket channel.
2. The initiating client requests the peer's public identity key and begins the handshake by sending its signed ephemeral key.
3. The peer responds with its own signed ephemeral key.
4. Both clients compute a shared secret using X25519 and derive two symmetric keys via HKDF (one for sending, one for receiving).
5. Encrypted messages are transmitted through the server as opaque ciphertexts.

Each step of this process is designed to ensure one or more core security properties:

- **Authentication** — guaranteed by the use of persistent Ed25519 identity keys and signature verification.
- **Confidentiality** — ensured by AES-256-GCM encryption of all message payloads.
- **Integrity** — provided by AES-GCM authentication tags and TLS channel protections.
- **Forward secrecy** — achieved by using ephemeral keys for each session.



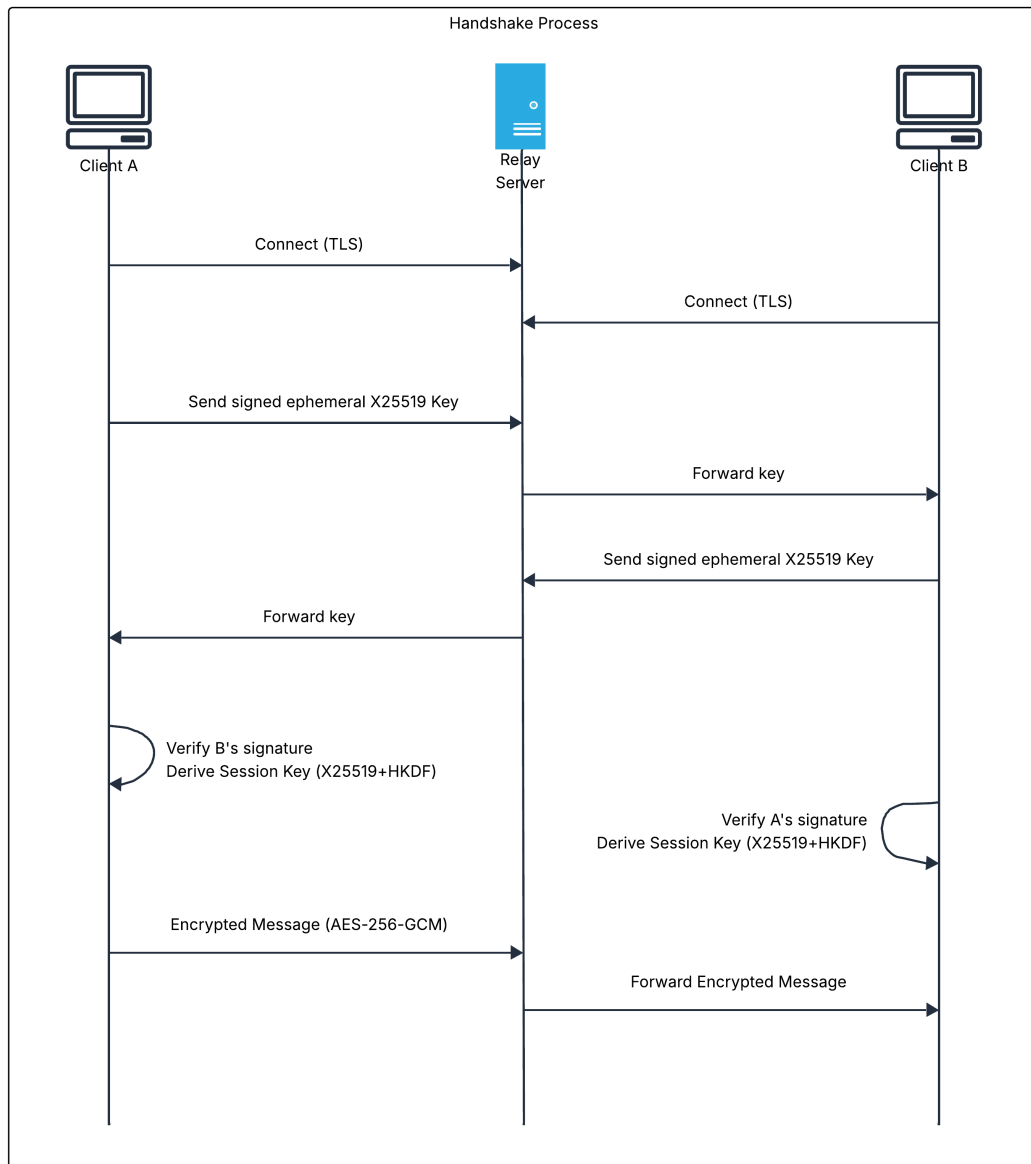


Figure 2: Handshake and key exchange: authenticated ephemeral key exchange using Ed25519 signatures and X25519 Diffie–Hellman, producing session keys via HKDF.

## 4 Design and Implementation

This chapter describes in detail how the cryptographic components, communication logic, and supporting libraries are integrated to achieve a secure end-to-end chat application. The focus is on the practical aspects of implementing the design while adhering to modern security best practices.

### 4.1 Key Management

Secure key management is the cornerstone of any end-to-end encrypted system. In our implementation, keys are generated and stored locally by each client, never leaving the device in unencrypted form.

- **Persistent Identity Key** Each user generates an Ed25519 key pair upon first launch of the client. The **private key** is stored locally in a secure format (not transmitted to the server or other clients), while the **public key** acts as the user's cryptographic identity. This persistent key allows peers to authenticate each other in future sessions, preventing impersonation attacks.
- **Ephemeral Key Exchange** For every new chat session, a fresh X25519 key pair is generated. Ephemeral keys provide **forward secrecy**: if a session key is ever compromised, past and future conversations remain secure because they use different ephemeral key material.
- **Signature Verification** Before an ephemeral public key is accepted, it is signed using the sender's persistent Ed25519 private key and verified by the recipient using the known persistent public key. This prevents man-in-the-middle (MITM) attacks by ensuring that the ephemeral key truly belongs to the claimed identity.

### 4.2 Session Key Derivation

Once both parties have exchanged and authenticated ephemeral public keys, they perform an X25519 elliptic-curve Diffie–Hellman (ECDH) computation to derive a shared secret.

- The shared secret is passed through an HKDF (HMAC-based Key Derivation Function) using SHA-256 as the hash function.
- From this HKDF output, **two distinct symmetric keys** are derived: one for sending messages and one for receiving messages. This key separation prevents subtle cryptographic vulnerabilities that could arise if the same key were reused in both directions.
- The initiator and responder deterministically derive *mirrored keys* so that each one's encryption key corresponds to the other's decryption key.

This approach ensures that session keys are:

- Unique for each session.
- Unlinkable between sessions.
- Resistant to key compromise (forward secrecy).

## 4.3 Message Encryption

### 4.3.1 AES-GCM

The symmetric encryption layer uses AES-256 in **Galois/Counter Mode** (AES-GCM), a modern authenticated encryption with associated data (AEAD) scheme.

AES-GCM provides:

- **Confidentiality** — The plaintext message content is encrypted using a per-message nonce.
- **Integrity and Authenticity** — Each ciphertext includes a cryptographic authentication tag that is verified before decryption. If verification fails, the message is discarded.

GCM mode uses a combination of:

- A *nonce* — a unique value per message, preventing replay and ensuring that identical plaintexts encrypt to different ciphertexts.
- A *counter* — incremented for each message to derive new nonces in a predictable but non-repeating way.

This makes AES-GCM resistant to a variety of attacks, including:

- Ciphertext pattern analysis.
- Replay attacks.
- Bit-flipping attacks (due to authentication tag verification).

The algorithm is computationally efficient and hardware-accelerated on most modern processors, which is why it is also the default choice in protocols such as TLS 1.3 and IPsec.

In our application, AES-GCM operates entirely at the *application layer* above TLS, meaning that messages remain encrypted even if the transport layer is compromised.

- AES-GCM (with associated data) ensures confidentiality + integrity.
- Nonces are derived from a per-message counter to guarantee uniqueness.

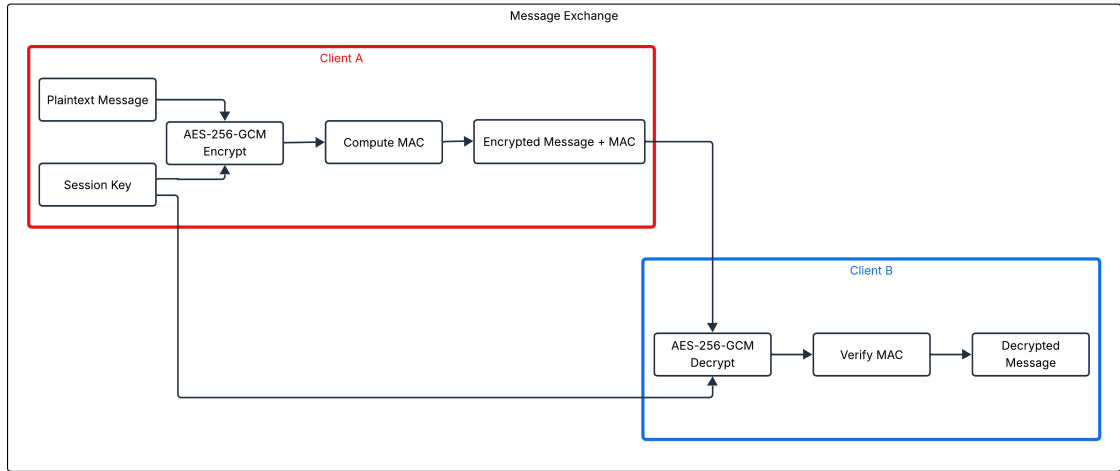


Figure 3: Message encryption workflow: plaintext is encrypted with AES-256-GCM using session keys, producing ciphertext with authentication tags that is sent over the TLS-protected WebSocket.

## 4.4 Broadcast Chat Functionality

Beyond one-to-one secure messaging, the system supports a **broadcast chat mode** in which multiple clients participate in a shared group conversation. This functionality builds on the same cryptographic foundation as the pairwise protocol, while introducing a group key management mechanism to ensure that only legitimate participants can access broadcast messages.

### 4.4.1 Joining a Broadcast

When a new client joins a broadcast session, it acts as the *committer*:

1. It requests the current membership list from the server.
2. It generates a fresh random **epoch secret** that defines the group's cryptographic state.
3. For each current member, it encrypts (wraps) the epoch secret with that member's public X25519 key.
4. It sends the updated membership list and wrapped secrets to the server in a **group\_commit** message.

Each member decrypts the wrapped secret using its private X25519 key and updates its local state to the new epoch.

### 4.4.2 Sending Broadcast Messages

To send a broadcast message:

1. The sender derives a per-sender AES-256-GCM key from:

- the epoch secret,
  - the group identifier,
  - the sender’s username (to ensure key separation between members).
2. A message counter serves as the nonce, guaranteeing uniqueness across the sender’s messages.
  3. The ciphertext is transmitted as a **group\_message** through the server, which relays it to all group members except the originator.

#### 4.4.3 Receiving Broadcast Messages

Upon receiving a broadcast message:

1. The recipient derives the same per-sender AES-GCM key.
2. Using the sender’s counter as the nonce, it decrypts the ciphertext.
3. Only members holding the current epoch secret are able to decrypt successfully.

#### 4.4.4 Leaving a Broadcast

When a member leaves:

1. A new epoch secret is generated by the committer.
2. The secret is wrapped for the remaining members, excluding the leaver.
3. This ensures that past messages remain confidential (backward secrecy) and the departing member cannot read future traffic (forward secrecy).

#### 4.4.5 Security Analysis

The broadcast protocol preserves the core guarantees of the one-to-one system:

- **Confidentiality:** Only members holding the current epoch secret can decrypt broadcast messages. The server never learns epoch secrets or derived keys.
- **Authentication and Integrity:** Messages are encrypted and authenticated with AES-256-GCM, and sender identity is tied to persistent Ed25519 keys.
- **Forward Secrecy:** Epoch rotation on membership changes prevents past or future compromise of group messages.
- **Resistance to Key/Nonce Collisions:** Per-sender derived keys and counters ensure uniqueness across group members, eliminating nonce reuse.
- **Minimal Trust in Server:** The server acts solely as a relay for opaque ciphertext and wrapped keys; compromise of the server reveals no plaintext.

This design resembles modern group messaging protocols such as MLS (Messaging Layer Security), though in a simplified form. It provides a balance between efficiency and strong security guarantees suitable for a relay-based broadcast architecture.

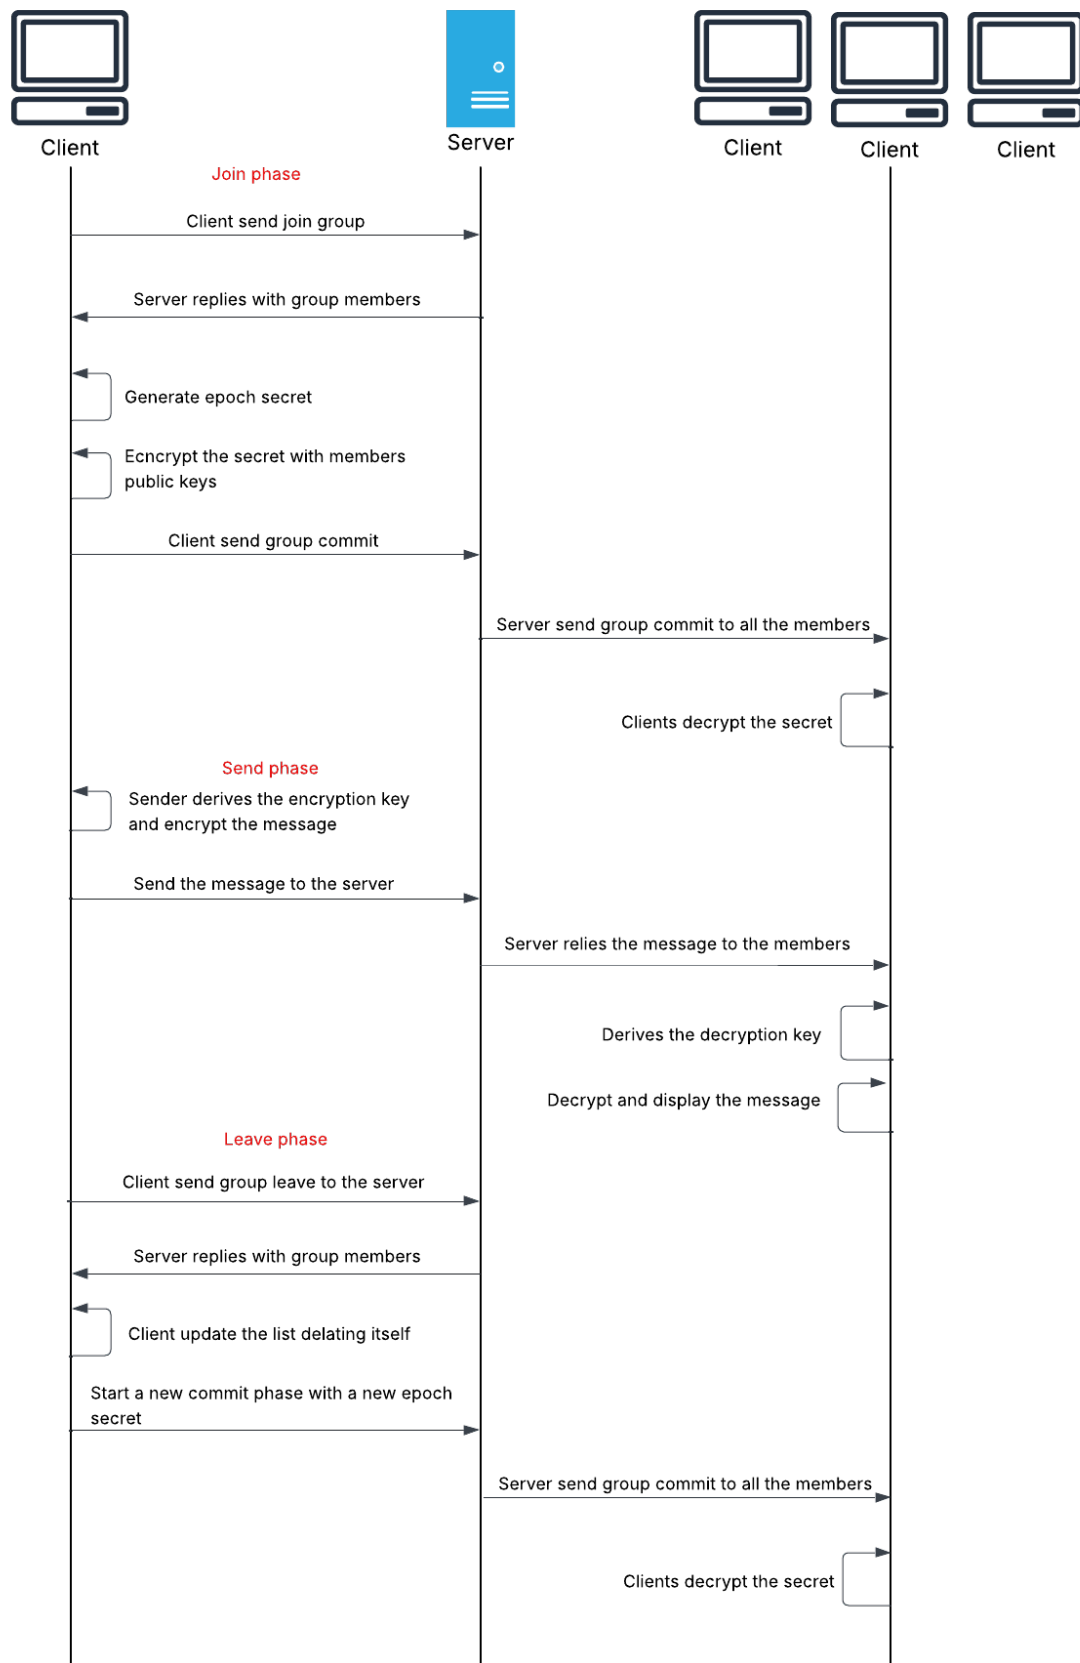


Figure 4: Key exchange, sending, receiving, joining and leaving the broadcast chat mechanism

## 4.5 Libraries Used

We selected Python libraries that provide secure, well-maintained, and high-level abstractions for cryptographic and networking tasks:

- **cryptography** — Core cryptographic primitives: Ed25519 (signatures), X25519 (ephemeral key exchange), AES-GCM (authenticated encryption), and HKDF-SHA256 (key derivation). This library underpins the security of our protocol implementation.
- **websockets** — Enables asynchronous client-server communication over WebSockets, with built-in support for TLS encryption.
- **asyncio** — Python’s event-driven concurrency framework, used to handle multiple clients, maintain responsive chat sessions, and manage asynchronous key exchanges.
- **tkinter** — Provides a simple graphical interface for interaction, user prompts, and session management.
- Standard library modules (**os**, **base64**, **hashlib**, **json**) — Support utilities for secure random generation, encoding/decoding keys and ciphertexts, hashing, and structured data exchange.

## 4.6 Code Structure

To maintain modularity and clarity, the project is organised into separate components:

- **client.py** — Handles the user interface, handshake protocol, and message encryption/decryption.
- **server.py** — Implements the WebSocket relay server for user discovery and message routing.
- **common/crypto\_utils.py** — Encapsulates all cryptographic operations, including key generation, signing, verification, ECDH, HKDF derivation, and AES-GCM encryption/decryption.

This separation of concerns simplifies debugging, testing, and future extensions to the system.

## 4.7 Interface

When the client is launched from the terminal, the secure messaging application opens and prompts the user to enter a username. Once the username is submitted, the main interface is displayed to the user.

The interface is divided into three main sections. The top section displays the user’s identity, including the username and fingerprint, and the status of the currently selected chat. On the left side, the interface displays all the active clients and the broadcast chat, where a user can start or end client-to-client chats, or join and leave the broadcast chat by selecting the desired entry and using the corresponding buttons. Finally, the

right section of the interface contains the selected chat window, where all incoming and outgoing messages are displayed. At the bottom of this section, a text input field allows users to write new messages and send them by clicking the send button.

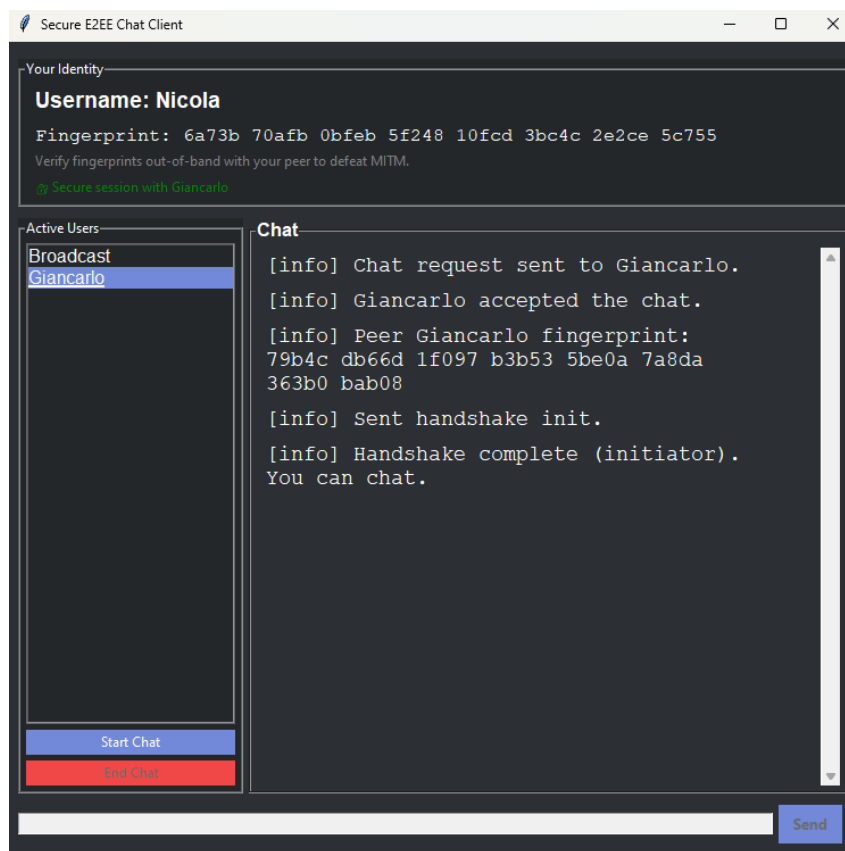


Figure 5: Chat interface



## 5 Evaluation and Conclusions

In this chapter, we report the results obtained by analysing network traffic with Wireshark to demonstrate the security of the communications carried out with the application we developed.

We have identified the packets that correspond to the main functions of our application, including the handshake, starting a chat with another client, joining the broadcast chat, and finally sending messages in both chats. From the highlighted packets, it is not possible to retrieve meaningful plaintext about clients or payloads. To identify these packets we applied Wireshark filters based on packet length and by performing the operations at distinct instants, which allowed us to map observed records to application actions.

### 5.1 Initial Handshakes

To evaluate the security of the system, we conducted a packet capture while three clients were connected to the server.

This setup allowed testing both one-to-one secure chats and the group broadcast functionality, while simultaneously monitoring the exchanged traffic with Wireshark.

For clarity of presentation, we filtered out packets shorter than 77 bytes. These packets are still part of the TLS communication, but in practice they mostly correspond to acknowledgements, retransmissions, or very small encrypted TLS records that are not relevant for illustrating the main application-level operations.

Filtering them improves visualisation and focuses the analysis on packets that represent significant cryptographic events (handshakes, key updates, and application ciphertext carrying chat messages). In the raw capture the total packet count is higher because TCP and TLS produce auxiliary traffic to ensure reliable and secure delivery.

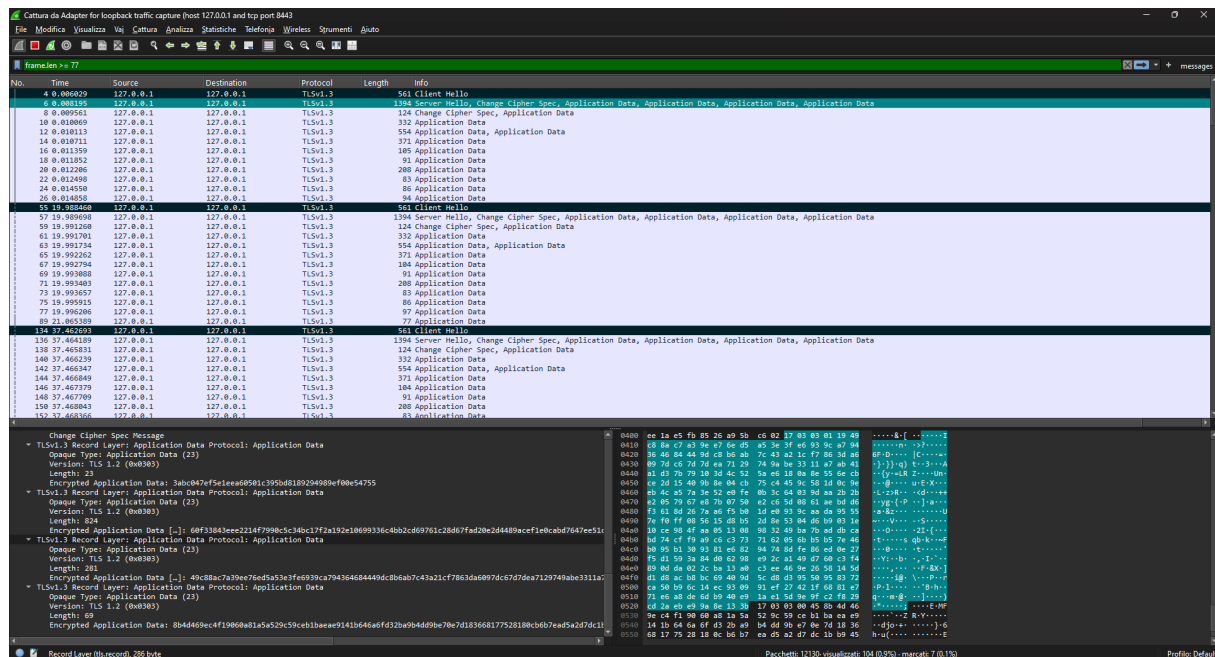


Figure 6: TLS handshake packets exchanged between clients and server

As shown in Figure 6, each client initiates communication with the server through a standard TLS handshake. Every session begins with a *ClientHello* message, followed by the server’s response (*ServerHello*) and the subsequent key exchange. While the protocol-level items (message types, lengths, and timing) are visible to an observer, the negotiated cryptographic secrets and session keys remain opaque. This confirms that no attacker observing the traffic could reconstruct the keys or decrypt subsequent communication.

## 5.2 Private (client-to-client) messages

Figure 7 highlights the packets corresponding to sending a private message from one client to another. The first highlighted packet in the capture is the sender’s TLS record directed to the server; the subsequent packet is the server routing the encrypted TLS record to the recipient.

By inspecting these records with Wireshark we can verify the following without exposing plaintext:

- The application-level payloads are carried inside TLS application records and are encrypted end-to-end with session keys negotiated inside TLS and with additional application-level encryption (X25519-derived ephemeral shared secret + AES-GCM) for the message bodies.
- Message integrity and authenticity are provided: AES-GCM provides authenticated encryption for ciphertexts and Ed25519 signatures are used at the application layer to bind sender identity to the message, preventing origin forgery.
- Replay protection is enforced: messages use unique nonces and associated data so that a recorder/replayer cannot successfully replay a previously captured ciphertext without detection.

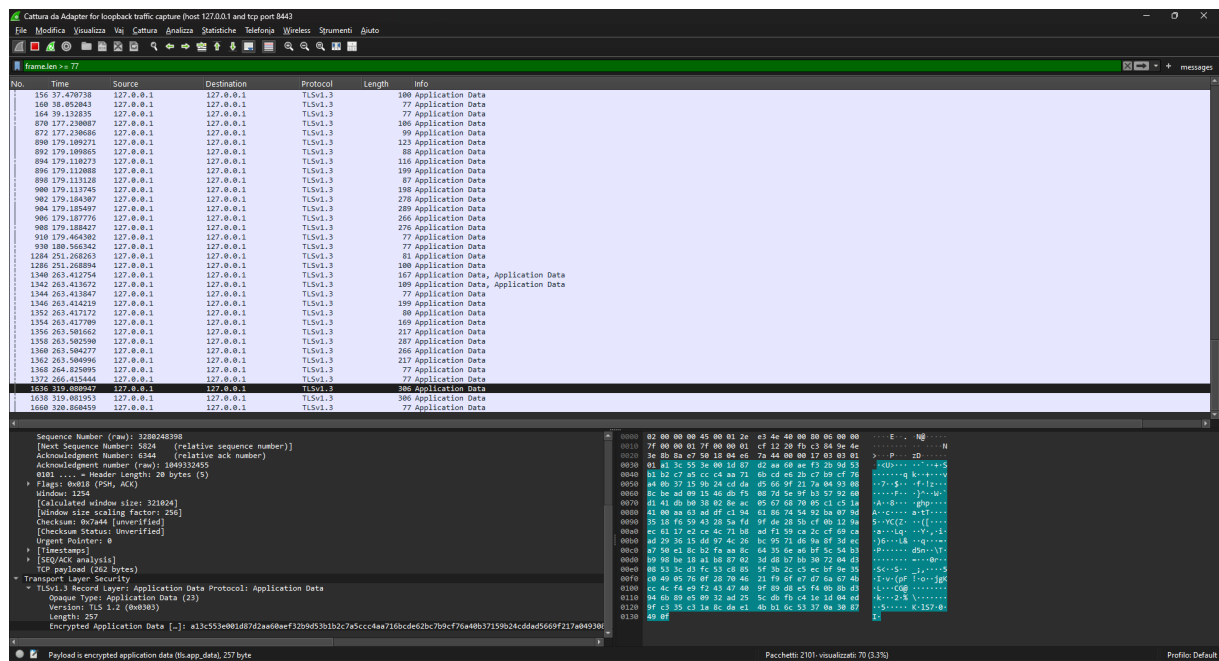


Figure 7: Packets observed when sending a private message (sender → server → recipient).

Because the ciphertexts, nonces, and associated authentication tags are all transmitted within encrypted TLS records, an external observer cannot retrieve message contents or replay them successfully without access to the appropriate keys.

## 5.2.1 Broadcast (group) messages and group key updates

Figure 8 shows the packets exchanged when a user sends a message in the broadcast chat. In the capture run there were three clients active, two of which were joined to the broadcast group. The server forwards the group ciphertext only to current group members; the packets illustrate that the server routes the encrypted application records to the relevant clients but does not expose plaintext content.

Importantly, the group mechanism implements per-epoch group keys with dynamic membership. When a member leaves the group, a new epoch key is distributed (encrypted for the remaining members) so that departed clients cannot decrypt future group messages. The capture shows key-update events (short encrypted records that differ from ordinary application ciphertexts) and subsequent messages protected under the new epoch key.

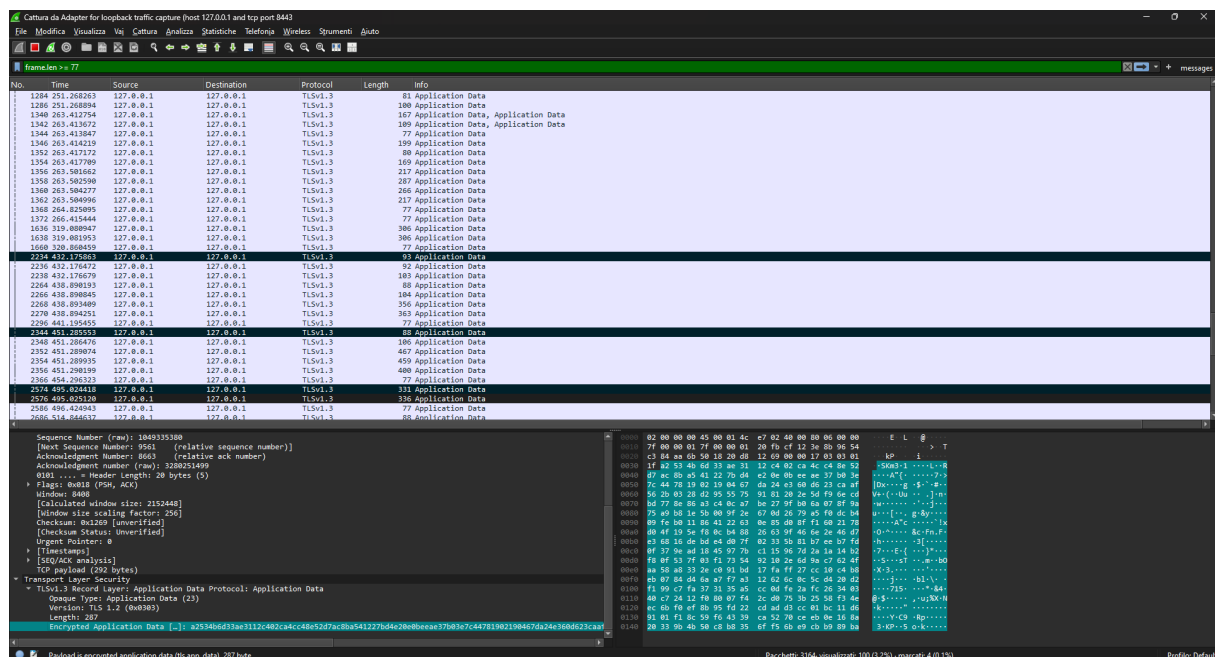


Figure 8: Encrypted broadcast messages forwarded by the server only to current group members; key-update events occur on membership changes.

## 5.3 What Wireshark reveals and what it does not

From packet inspection we can observe:

- Protocol-level metadata: TLS record boundaries, message types (e.g., ClientHello), packet lengths, source/destination IPs and ports, timing and packet order.
- The presence of application activity and roughly when key updates occur, since those events change record sizes and timing patterns.

What Wireshark (and a passive eavesdropper) cannot reveal:

- Application plaintext: all chat payloads remain encrypted and integrity-protected by AES-GCM and by TLS.
- Long-term private keys: Ed25519 private keys and X25519 ephemeral private keys are never transmitted on the wire.
- The exact session keys or HKDF outputs used to derive message keys.

## 5.4 Limitations

The tests above demonstrate the confidentiality and integrity properties of the implementation, but there are practical limitations to our current prototype and to this evaluation:

- Limited scalability: the prototype server and broadcast mechanism have not been stress-tested for large numbers of concurrent users or high message rates.
- No persistent message history or offline delivery: messages are only delivered to online recipients; persistent, encrypted storage is not implemented in the current prototype.
- Server-observable metadata: the server still learns routing metadata (which client sent a message and which client(s) received it) and can perform timing/traffic-analysis on communications.
- Use of self-signed TLS certificates: the server automatically generates a self-signed certificate at startup if none is present. While this guarantees that TLS is always available in a local development setting, it does not provide the authentication properties of a CA-issued certificate and will trigger warnings in standard clients.

## 5.5 Possible Improvements

We identify several practical enhancements and further tests that would strengthen the system:

- GUI improvements (bubble-style messages, user info, session status) for usability and clearer session indicators.
- Persistent, client-side encrypted chat history with forward secrecy for stored messages.
- Stronger forward secrecy across multiple sessions and a defined ratchet mechanism for long-lived conversations.
- Scalability and performance testing (load tests).
- Minimising server metadata leakage.
- Deployment with CA-signed TLS certificates to ensure strong server authentication and eliminate client-side trust warnings.

## 5.6 Conclusion

The application implements the core principles of secure messaging: identity verification (Ed25519), ephemeral session key negotiation (X25519), authenticated encryption (AES-GCM), and transport security (TLS). Our packet captures and Wireshark analysis confirm that message contents are not recoverable from passive network traces and that group key updates prevent departed members from decrypting subsequent group messages.

While this prototype is suitable for demonstration and teaching purposes, further work is required to harden it for production use: add persistent encrypted storage, perform formal verification and penetration testing, and reduce server-observable metadata. Nevertheless, the tests performed provide practical evidence that the implemented cryptographic primitives and protocols protect confidentiality and integrity as intended.